

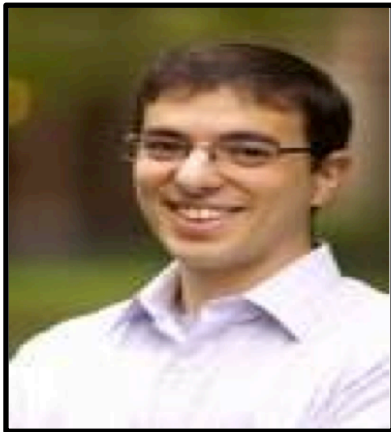
Lazy Evaluation of Transactions in Database Systems

Jose M. Faleiro◆, Alexander Thomson★,
Daniel J. Abadi◆

◆Yale University, ★Google




Paper “X” looks relevant,
can you add it to related
work?



Just read through the whole
thing and added to related work.
Our work is more ACIDic





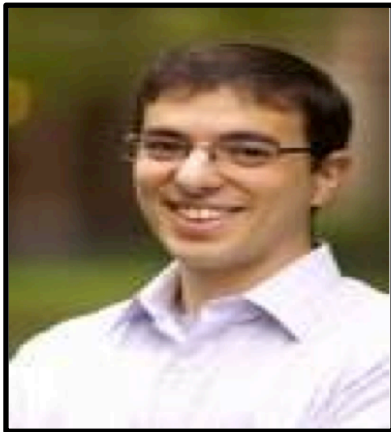
In reality...

- Read only abstract to give a sensible response
- Keep “sticky” note on desk to later read paper and write related work.



Read paper X

Add to
related work



Read paper X

Add to
related work

Read paper Y

Add to
related work

...

Read paper Z

Add to
related work

Hey, can you check-in your additions to related work? I can't see your changes...



Read paper X

Add to
related work

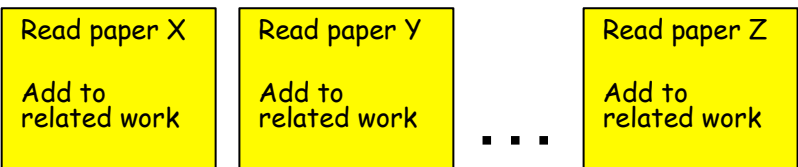
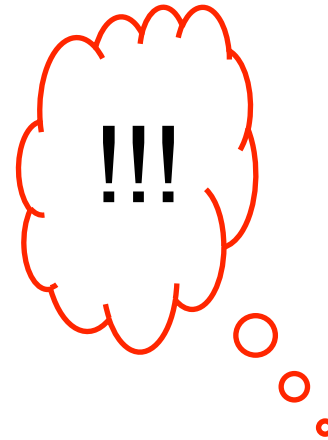
Read paper Y

Add to
related work

...

Read paper Z

Add to
related work



- Previously said I did all the work. Can't feign ignorance.
- Use information on stickies to read papers and finish writing related work.



Read paper X

Add to
related work

Read paper Y

Add to
related work

...

Read paper Z

Add to
related work

Lazy Database Systems

- Promise to commit/abort while only partially executing txns
- Keep promises while maintaining ACID and serializability

Lying Considered ~~Harmful~~ Useful

- Flexibility to execute transactions when most favorable
 - Cache Locality
 - Load Balancing
 - Contention Reduction

Can't Keep Lying Forever



Must satisfy “external”
reads

Talk Outline

- **Lazy Database Design**
- Benefits of Lazy Execution
- Experimental Evaluation
- Conclusion

Customer Orders Product

Update product
inventory

Compute
discounted price

Update customer
bill

Customer Orders Product

Update product
inventory

Compute
discounted price

Update customer
bill

Read Product record

Product count -= 1

If Product count < 0:

ABORT()

Customer Orders Product

Update product
inventory

**Compute
discounted price**

Update customer
bill

```
Read Product record
```

```
Product count -= 1
```

```
If Product count < 0:
```

```
    ABORT()
```

```
Read daily discount
```

```
Compute discounted price
```

Customer Orders Product

Update product
inventory

Compute
discounted price

Update customer
bill

```
Read Product record  
  
Product count -= 1  
  
If Product count < 0:  
    ABORT()  
  
Read daily discount  
  
Compute discounted price
```

Read customer's monthly
bill record

Customer's monthly bill +=
Discounted price

COMMIT()

We Ask the Question

Can we:
return commit/abort
promise ...

without executing to
completion...

while maintaining
ACID and
serializability?

```
Read Product record
```

```
Product count -= 1
```

```
If Product count < 0:
```

```
    ABORT()
```

```
Read daily discount
```

```
Compute discounted price
```

```
Read customer's monthly  
bill record
```

```
Customer's monthly bill +=  
    Discounted price
```

```
COMMIT()
```

Our Solution

- Split up transactions into two parts
 - Give users the illusion that execution is atomic
- First part produces commit/abort promise
 - Executes immediately
- Second part does everything else
 - Execution is deferred

How to Split a Transaction?

Read Product record

Product count -= 1

If Product count < 0:

ABORT()

Read daily discount

Compute discounted price

Read customer's monthly
bill record

Customer's monthly bill +=
Discounted price

COMMIT()

Splitting a Transaction

Abort due to
txn logic

Read Product record

Product count -= 1

If Product count < 0:

ABORT()

Read daily discount

Compute discounted price

Read customer's monthly
bill record

Customer's monthly bill +=
Discounted price

COMMIT()

Splitting a Transaction

Read Product record

Product count -= 1

If Product count < 0:

ABORT()

Read daily discount

Compute discounted price

Read customer's monthly
bill record

Customer's monthly bill +=
Discounted price

COMMIT()

Return
commit
promise



Splitting a Transaction

Read Product record

Product count -= 1

If Product count < 0:

ABORT()

Read daily discount

Compute discounted price

Read customer's monthly
bill record

Customer's monthly bill +=
Discounted price

COMMIT()

Splitting a Transaction

Read Product record

Product count -= 1

If Product count < 0:

ABORT()

COMMIT_PROMISE()

Read daily discount

Compute discounted price

Read customer's monthly
bill record

Customer's monthly bill +=
Product price

Splitting a Transaction

Execute immediately



```
graph TD; A[Execute immediately] --> B[Read Product record<br/>Product count -= 1<br/>If Product count < 0:<br/>  ABORT()<br/>  COMMIT_PROMISE()]; C[Defer execution] --> D[Read daily discount<br/>Compute discounted price<br/>Read customer's monthly bill record<br/>Customer's monthly bill += Product price];
```

Read Product record

Product count -= 1

If Product count < 0:

ABORT()

COMMIT_PROMISE()

Read daily discount

Compute discounted price

Read customer's monthly
bill record

Customer's monthly bill +=
Product price


Defer execution

Dealing with Deferred Logic

- Don't write out actual record values
- ***Blindly*** insert placeholders corresponding to unexecuted deferred logic
 - Similar to stickies on my desk as reminders

Dealing with Deferred Logic

Monthly Bill

John	\$17.05
Mary	
⋮	
David	\$300

Read Tuesday's discount

Compute discounted price

Read customer's monthly bill record


Mary's monthly bill +=
Product price

Processing External Reads

- External read may depend on a deferred transaction
- Deferred transaction itself may have dependencies

Dealing with Deferred Logic

Monthly Bill

John	\$17.05
Mary	
⋮	
David	\$300


Read Tuesday's discount

Compute discounted price

Read customer's monthly bill record

Mary's monthly bill +=
Product price

Daily Discount

Monday	5%
Tuesday	
⋮	
Friday	15%

Increase Tuesday's discount by 10%

Dealing with Deferred Logic

Monthly Bill

John	\$17.05
Mary	<input type="text"/>
⋮	
David	\$300

Read Tuesday's discount

Compute discounted price

Read customer's monthly bill record

Mary's monthly bill += Product price

Daily Discount

Monday	5%
Tuesday	<input type="text"/>
⋮	
Friday	15%

Increase Tuesday's discount by 10%

Dealing with Deferred Logic

Deferred pieces of logic
are implicitly ordered

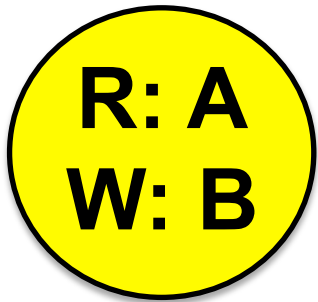
```
Read Tuesday's discount  
Compute discounted price  
Read customer's monthly  
bill record  
Mary's monthly bill +=  
    Product price
```



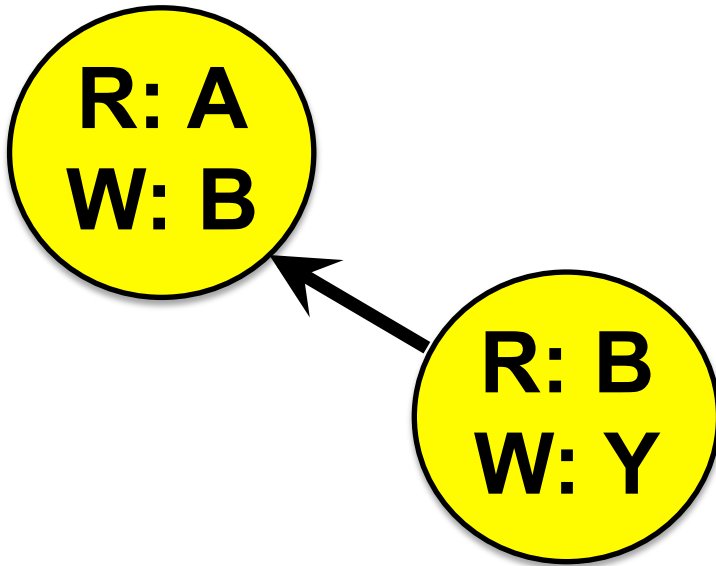
```
Increase Tuesday's  
discount by 10%
```


Dependency Graph of Transactions

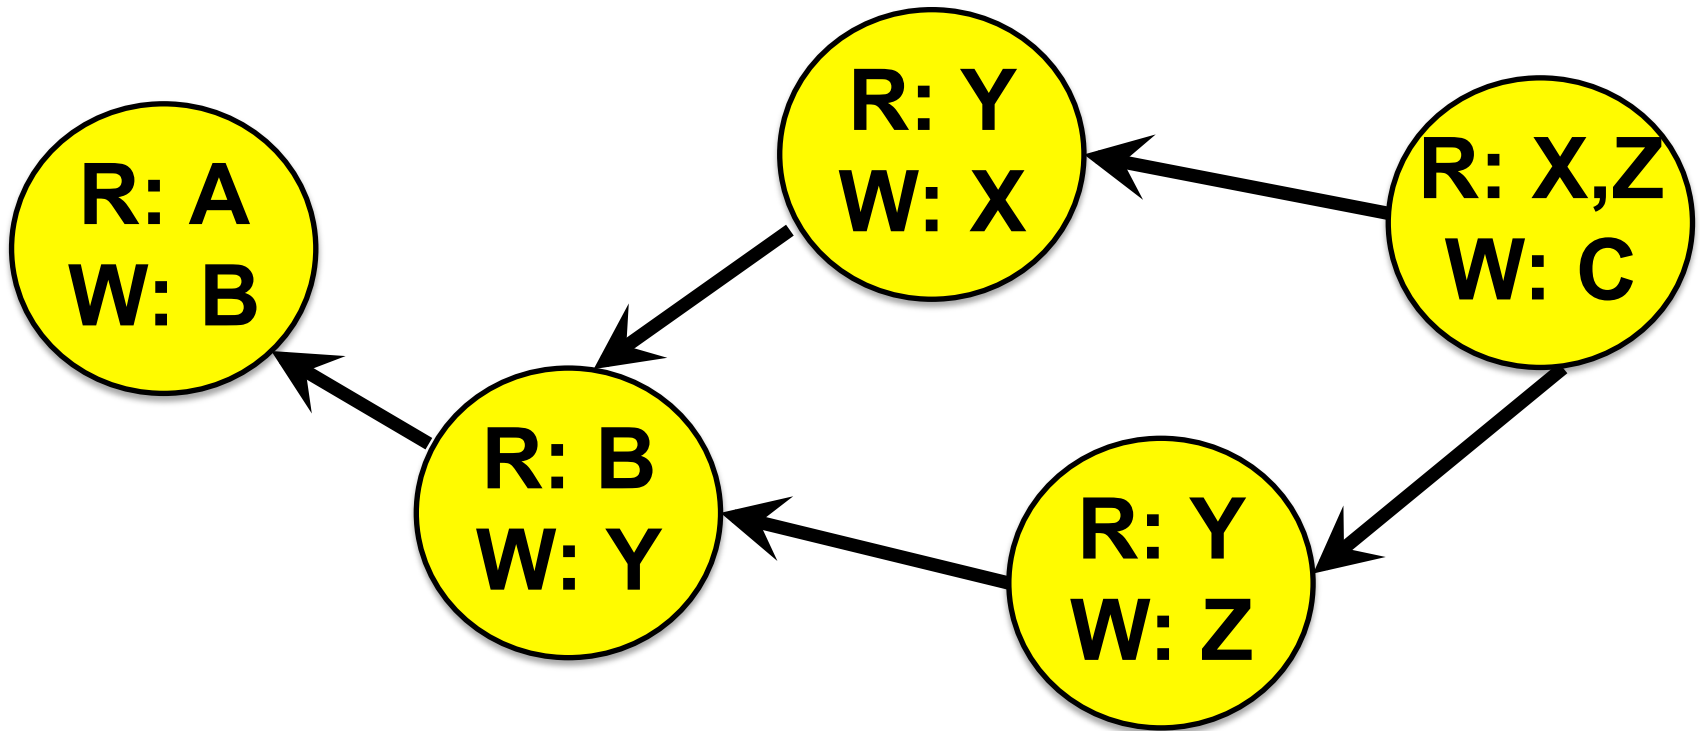
Dependency Graph of Transactions



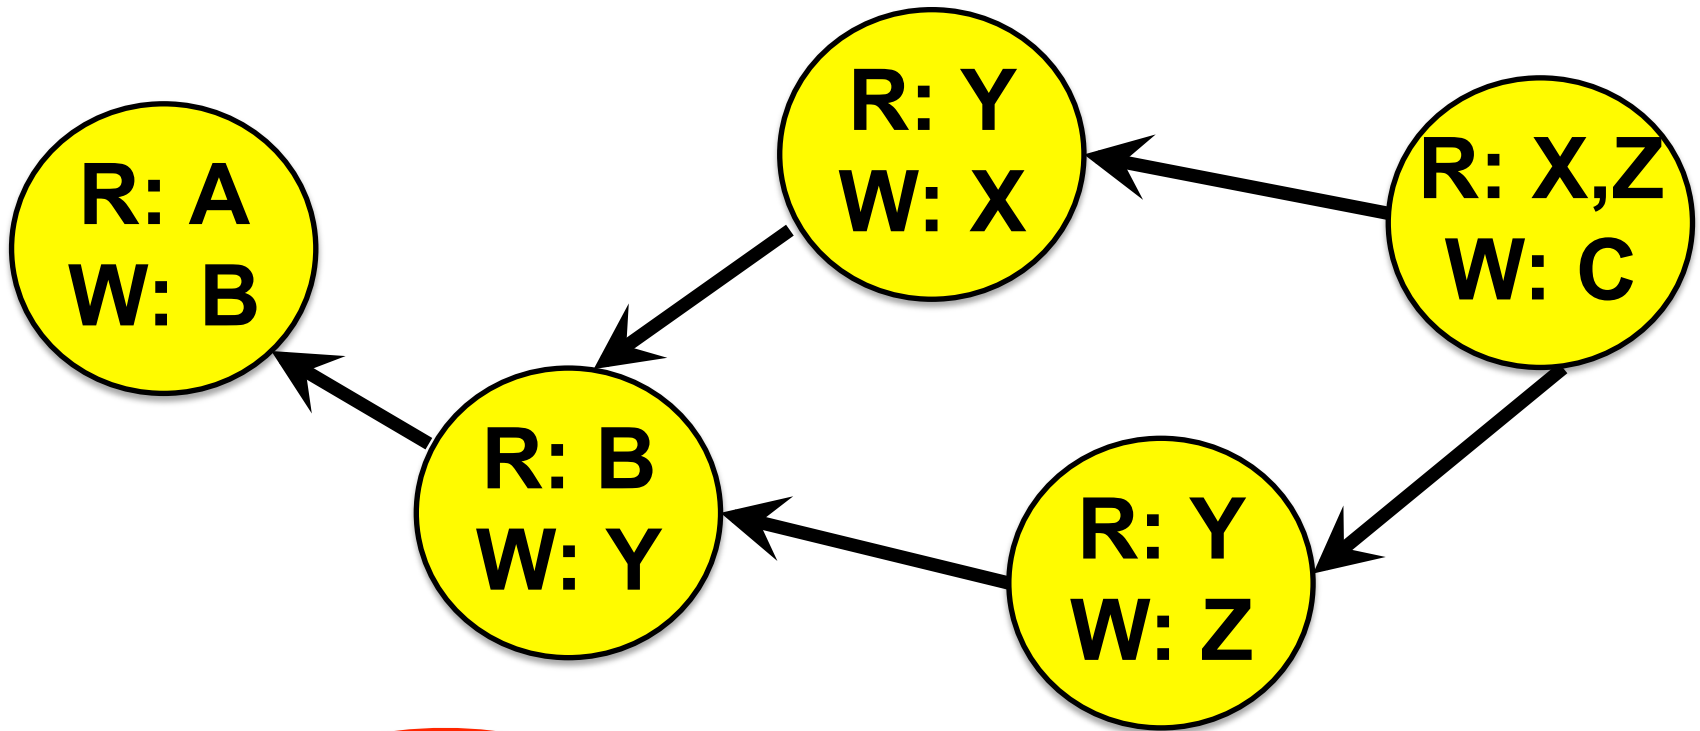
Dependency Graph of Transactions



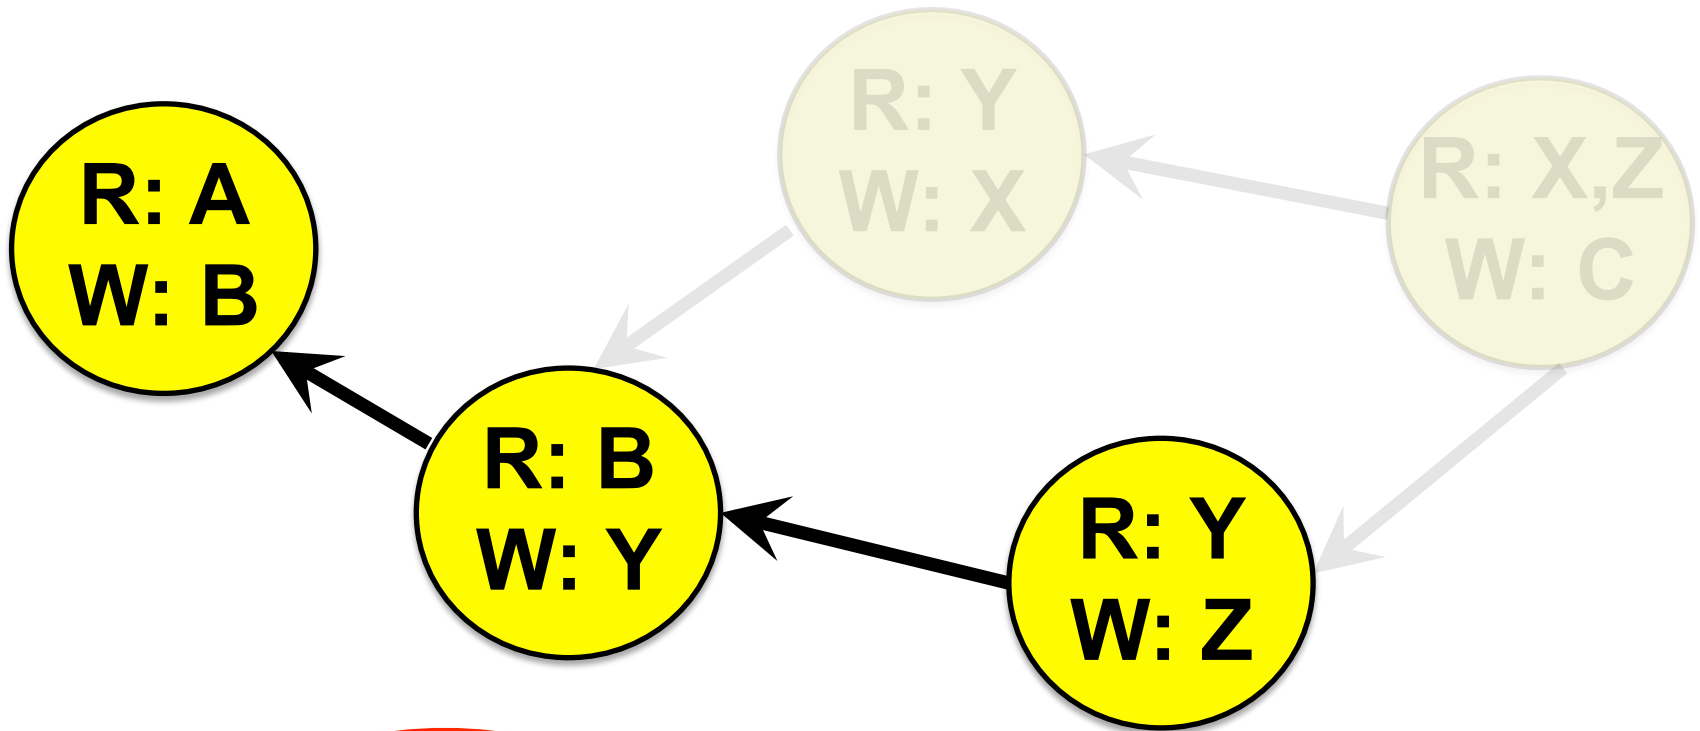
Dependency Graph of Transactions



Dependency Graph of Transactions



Dependency Graph of Transactions



Execute last writer
and its transitive
closure

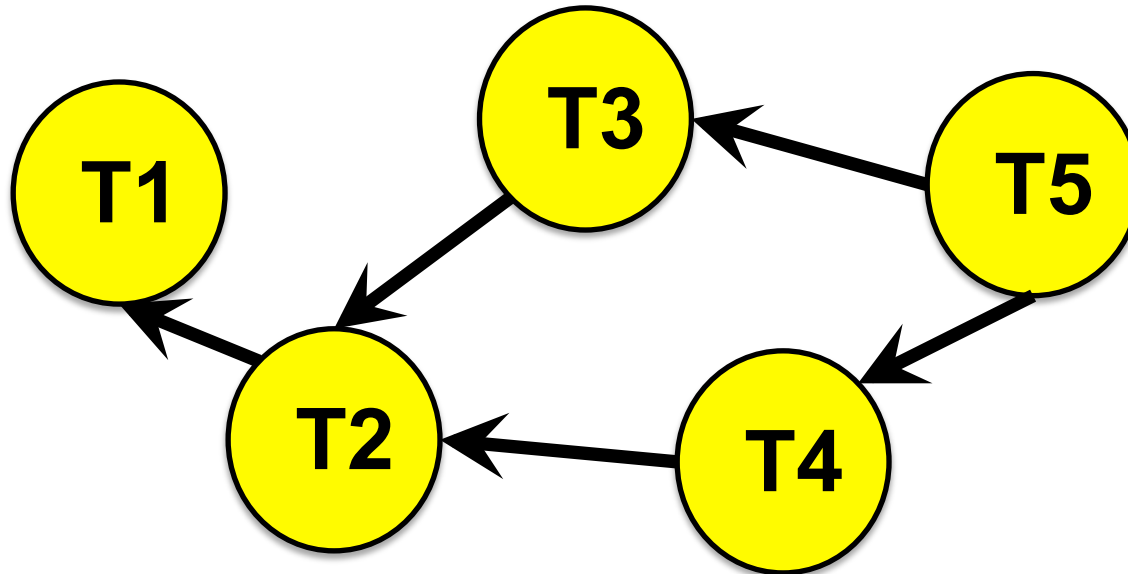
Large Transitive Closures

- The subset of the graph a record depends on may grow very large
 - High external read latency

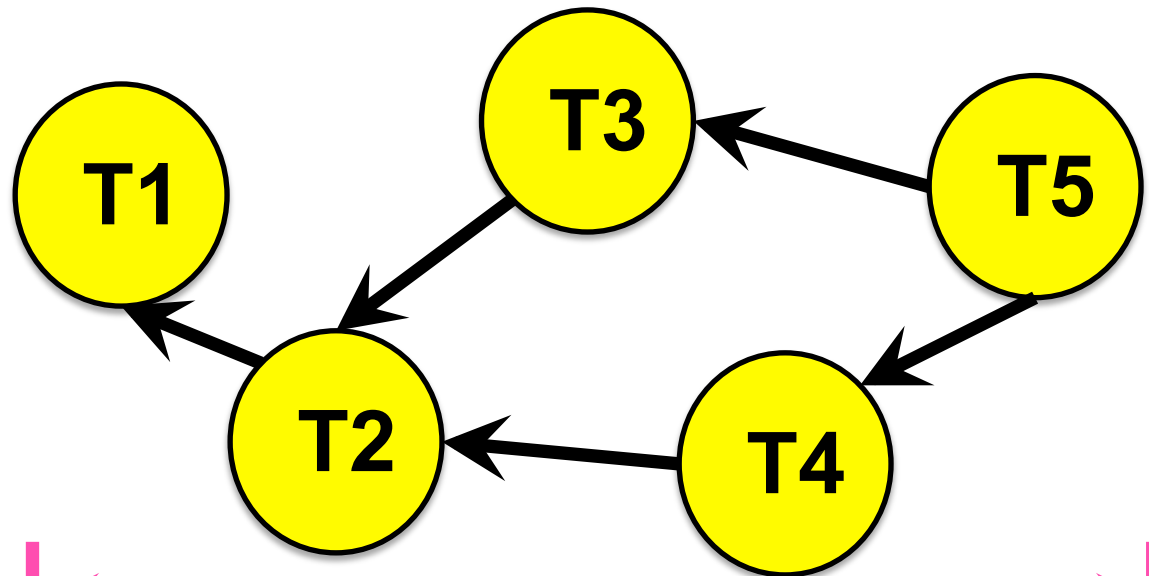
Large Transitive Closures

- The subset of the graph a record depends on may grow very large
 - High external read latency
 - **Incrementally process transactions in the background**

Incremental Transaction Processing




Incremental Transaction Processing



Bound the maximum
depth of the dependency
graph

Incremental Transaction Processing



Bound the maximum
depth of the dependency
graph

Incremental Transaction Processing

Tradeoff:

- Higher bound means larger batches (good for cache locality)
- Larger batches mean increased external read latency

Talk Outline

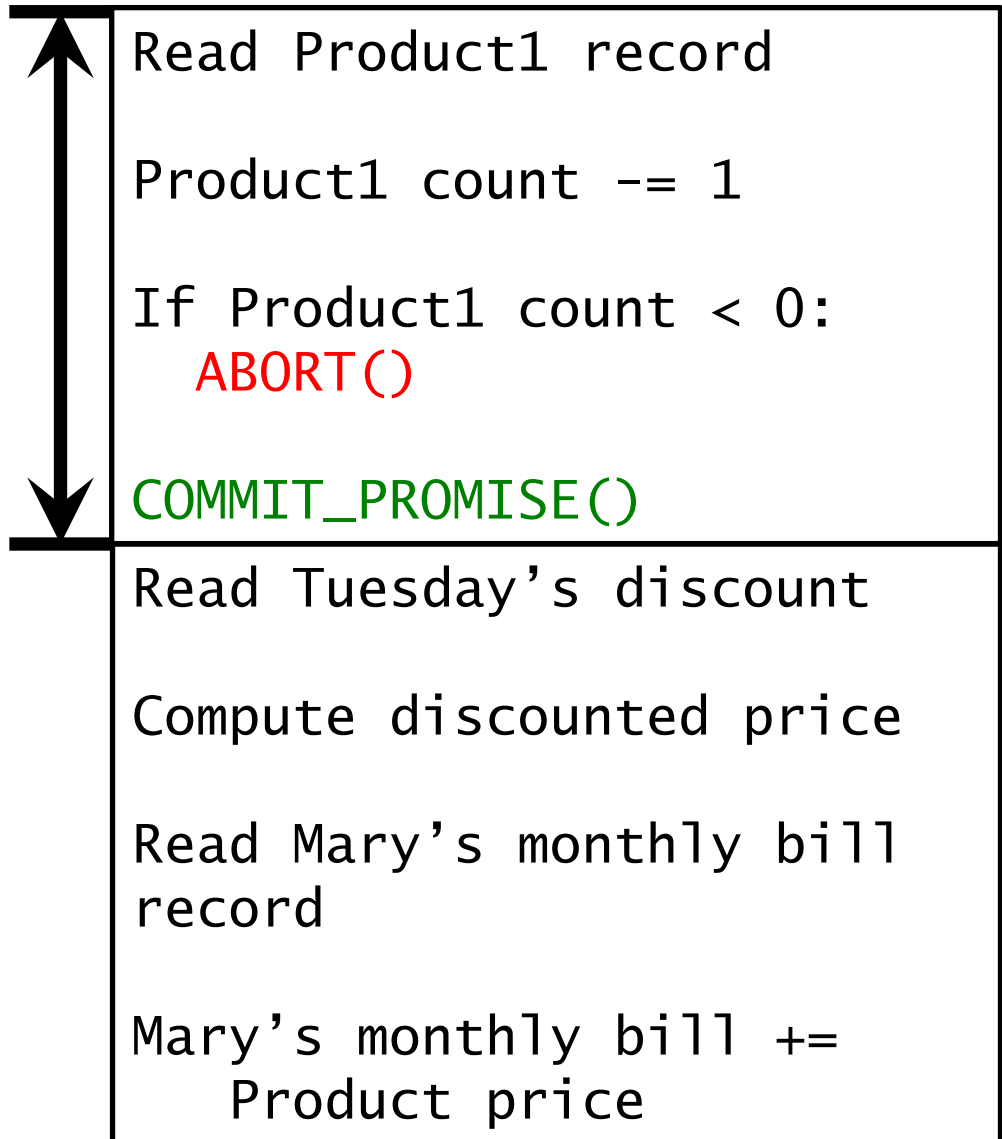
- Lazy Database Design
- **Benefits of Lazy Execution**
- Experimental Evaluation
- Conclusion

Benefits of Lazy Transactions

- Data Cache/Buffer Pool Locality
- Temporal Load Balancing
- Reduced Contention

Data Cache/Buffer Pool Locality

Execute
immediately



Data Cache/Buffer Pool Locality

```
Read Tuesday's discount  
Compute discounted price  
Read Mary's monthly bill  
record  
Mary's monthly bill +=  
    Product price
```

Deferred

Data Cache/Buffer Pool Locality

Read Mary's monthly bill
record

Mary's monthly bill +=
Product1 price

Insert a record into
Orders table

Deferred

Read Product2 record

Product2 count -= 1

If Product2 count < 0:

ABORT()

COMMIT_PROMISE()

Read Monday's discount

Compute discounted price

Read Mary's monthly bill
record

Mary's monthly bill +=
Product price

Data Cache/Buffer Pool Locality

Read Mary's monthly bill

re Read Monday's discount

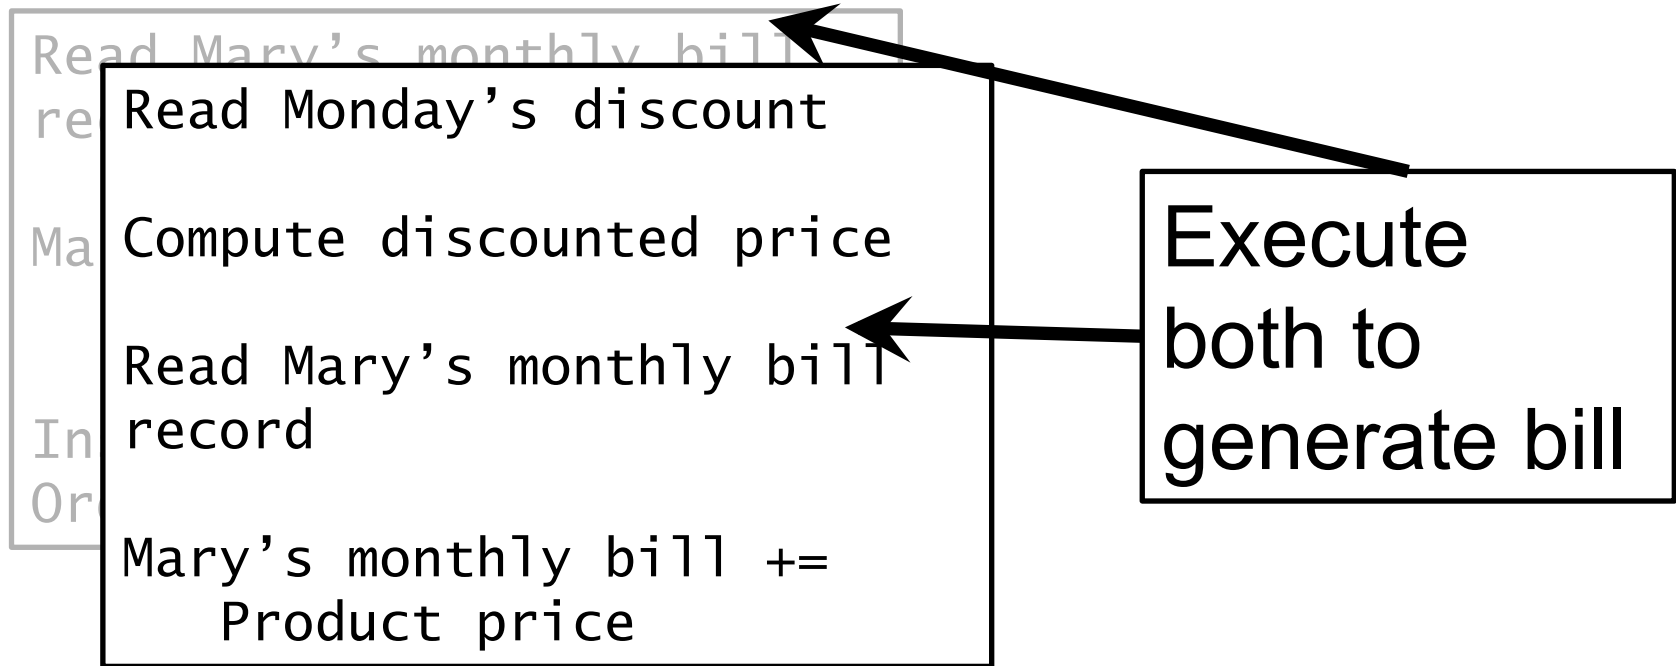
Ma Compute discounted price

In Read Mary's monthly bill
Or record

Mary's monthly bill +=
Product price

Deferred

Data Cache/Buffer Pool Locality



Deferred

Data Cache/Buffer Pool Locality

Read Monday's discount

Compute discounted price

Read Mary's monthly bill
record

Mary's monthly bill +=
Product2 price

Read Tuesday's discount

Compute discounted price

Read Mary's monthly bill
record

Mary's monthly bill +=
Product1 price

Deferred

Both update
the same bill

```
graph TD; A[Both update the same bill] --> B[Read Monday's discount<br/>Compute discounted price<br/>Read Mary's monthly bill record<br/>Mary's monthly bill += Product2 price]; A --> C[Read Tuesday's discount<br/>Compute discounted price<br/>Read Mary's monthly bill record<br/>Mary's monthly bill += Product1 price];
```

Data Cache/Buffer Pool Locality

Read Monday's discount

Compute discounted price

Read Mary's monthly bill
record

Mary's monthly bill +=
Product2 price

Read Tuesday's discount

Compute discounted price

Read Mary's monthly bill
record

Mary's monthly bill +=
Product1 price

Deferred

Bring Mary's bill
record into cache
just once



Temporal Load Balancing

Read Product record

Product count -= 1

If Product count < 0:

ABORT()

Read daily discount

Compute discounted price

Read Customer's monthly
bill record

Customer's monthly bill +=
Product price

COMMIT()

Read Product record

Product count -= 1

If Product count < 0:

ABORT()

COMMIT_PROMISE()

Read daily discount

Compute discounted price

Read Customer's monthly
bill record

Customer's monthly bill +=
Product price

Temporal Load Balancing

Read Product record

Product count -= 1

If Product count < 0:
ABORT()

Read daily discount

Compute discounted price

Read Customer's monthly
bill record

Customer's monthly bill +=
Product price

COMMIT()

Read Product record

Product count -= 1

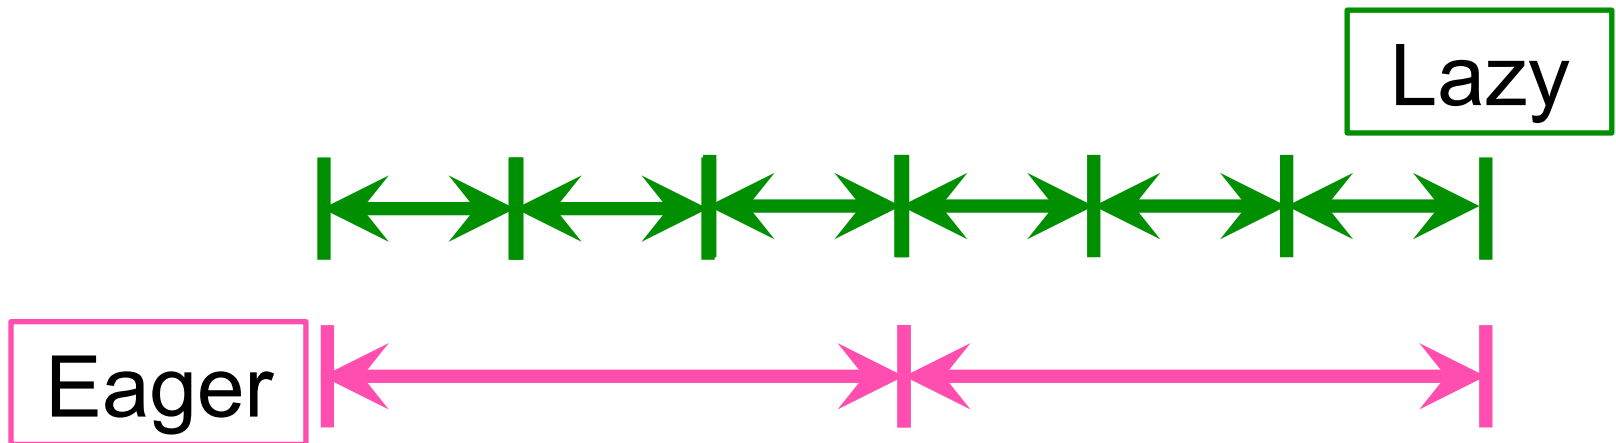
If Product count < 0:
ABORT()

COMMIT_PROMISE()

Arrows represent
time required to
immediately process
transactions

Temporal Load Balancing

More commit
decisions per
unit time



Reduced Contention

Popular item

Read **iThing** record

iThing count -= 1

If **iThing** count < 0:

ABORT()

Read daily discount

Compute discounted price

Read Fanboy1's monthly
bill record

Fanboy1's monthly bill +=
Discounted price

COMMIT()

Reduced Contention

Lock must
be held for
this period



```
Read iThing record
```

```
iThing count -= 1
```

```
If iThing count < 0:
```

```
    ABORT()
```

```
Read daily discount
```

```
Compute discounted price
```

```
Read Fanboy1's monthly  
bill record
```

```
Fanboy1's monthly bill +=  
    Discounted price
```

```
COMMIT()
```

Reduced Contention

Hold
contended
lock for less
time



```
Read iThing record  
  
iThing count -= 1  
  
If iThing count < 0:  
    ABORT()  
  
COMMIT_PROMISE()
```

```
Read daily discount  
  
Compute discounted price  
  
Read Fanboy1's monthly  
bill record  
  
Fanboy1's monthly bill +=  
    Discounted price
```

Talk Outline

- Lazy Database Design
- Benefits of Lazy Execution
- **Experimental Evaluation**
- Conclusion

Experimental Setup

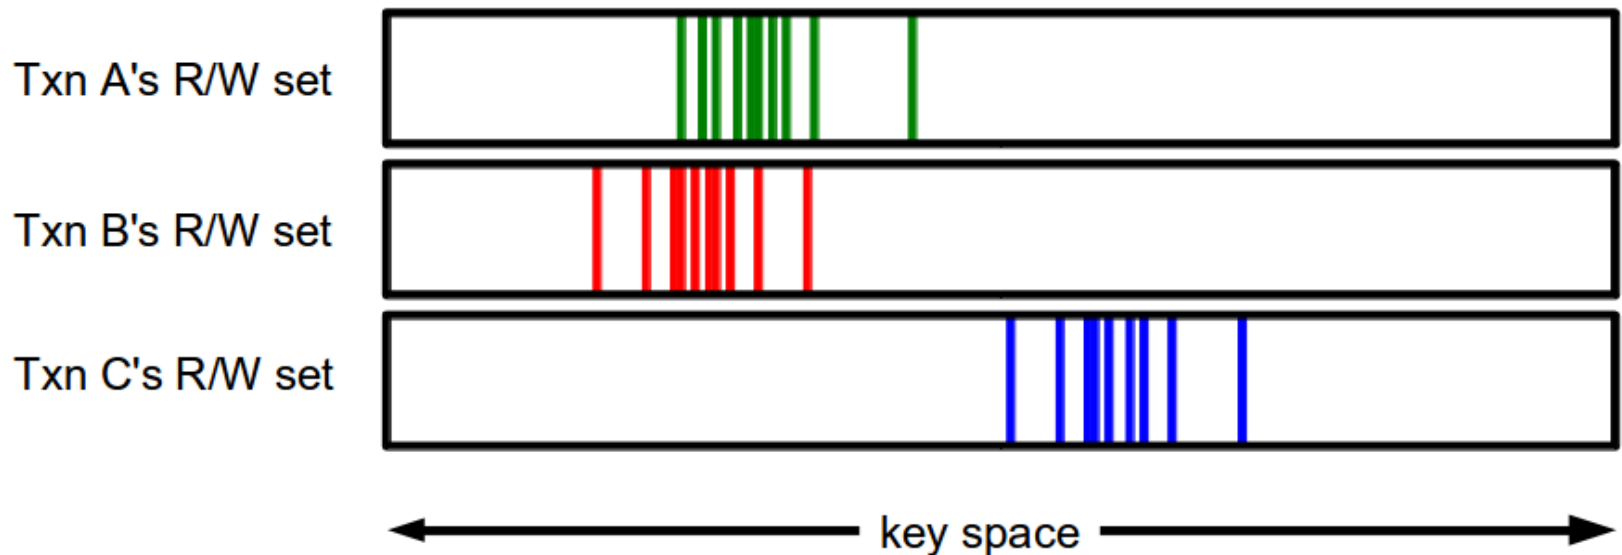
- All data in memory
- Single machine
- 8 cores
 - Lazy: 1 commit/abort + 7 deferred logic
 - Eager: 8 cores

Experimental Evaluation

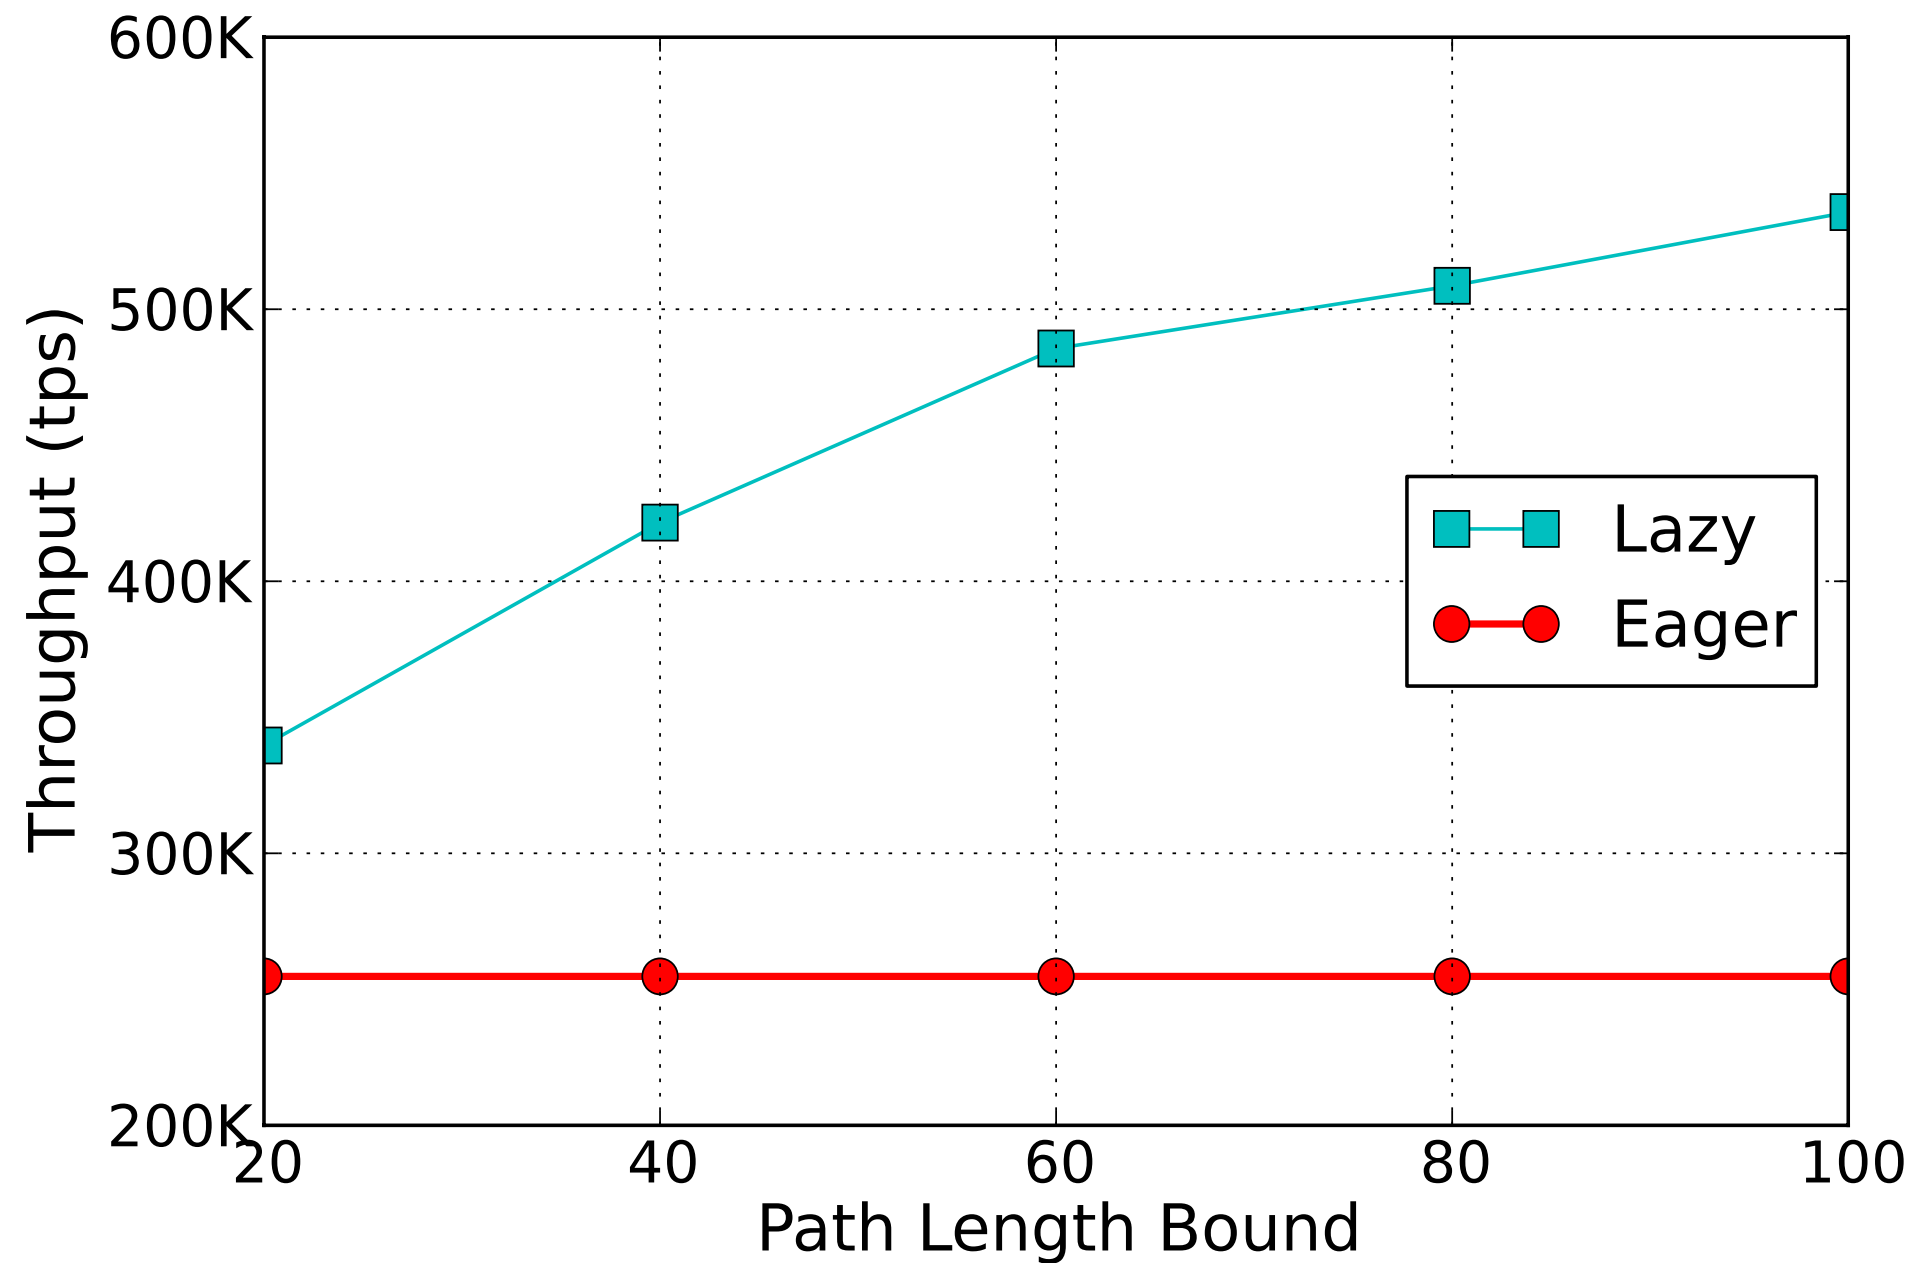
- Cache Locality
- External Read Latency
- Temporal Load Balancing
- Contention Reduction

Cache-Locality

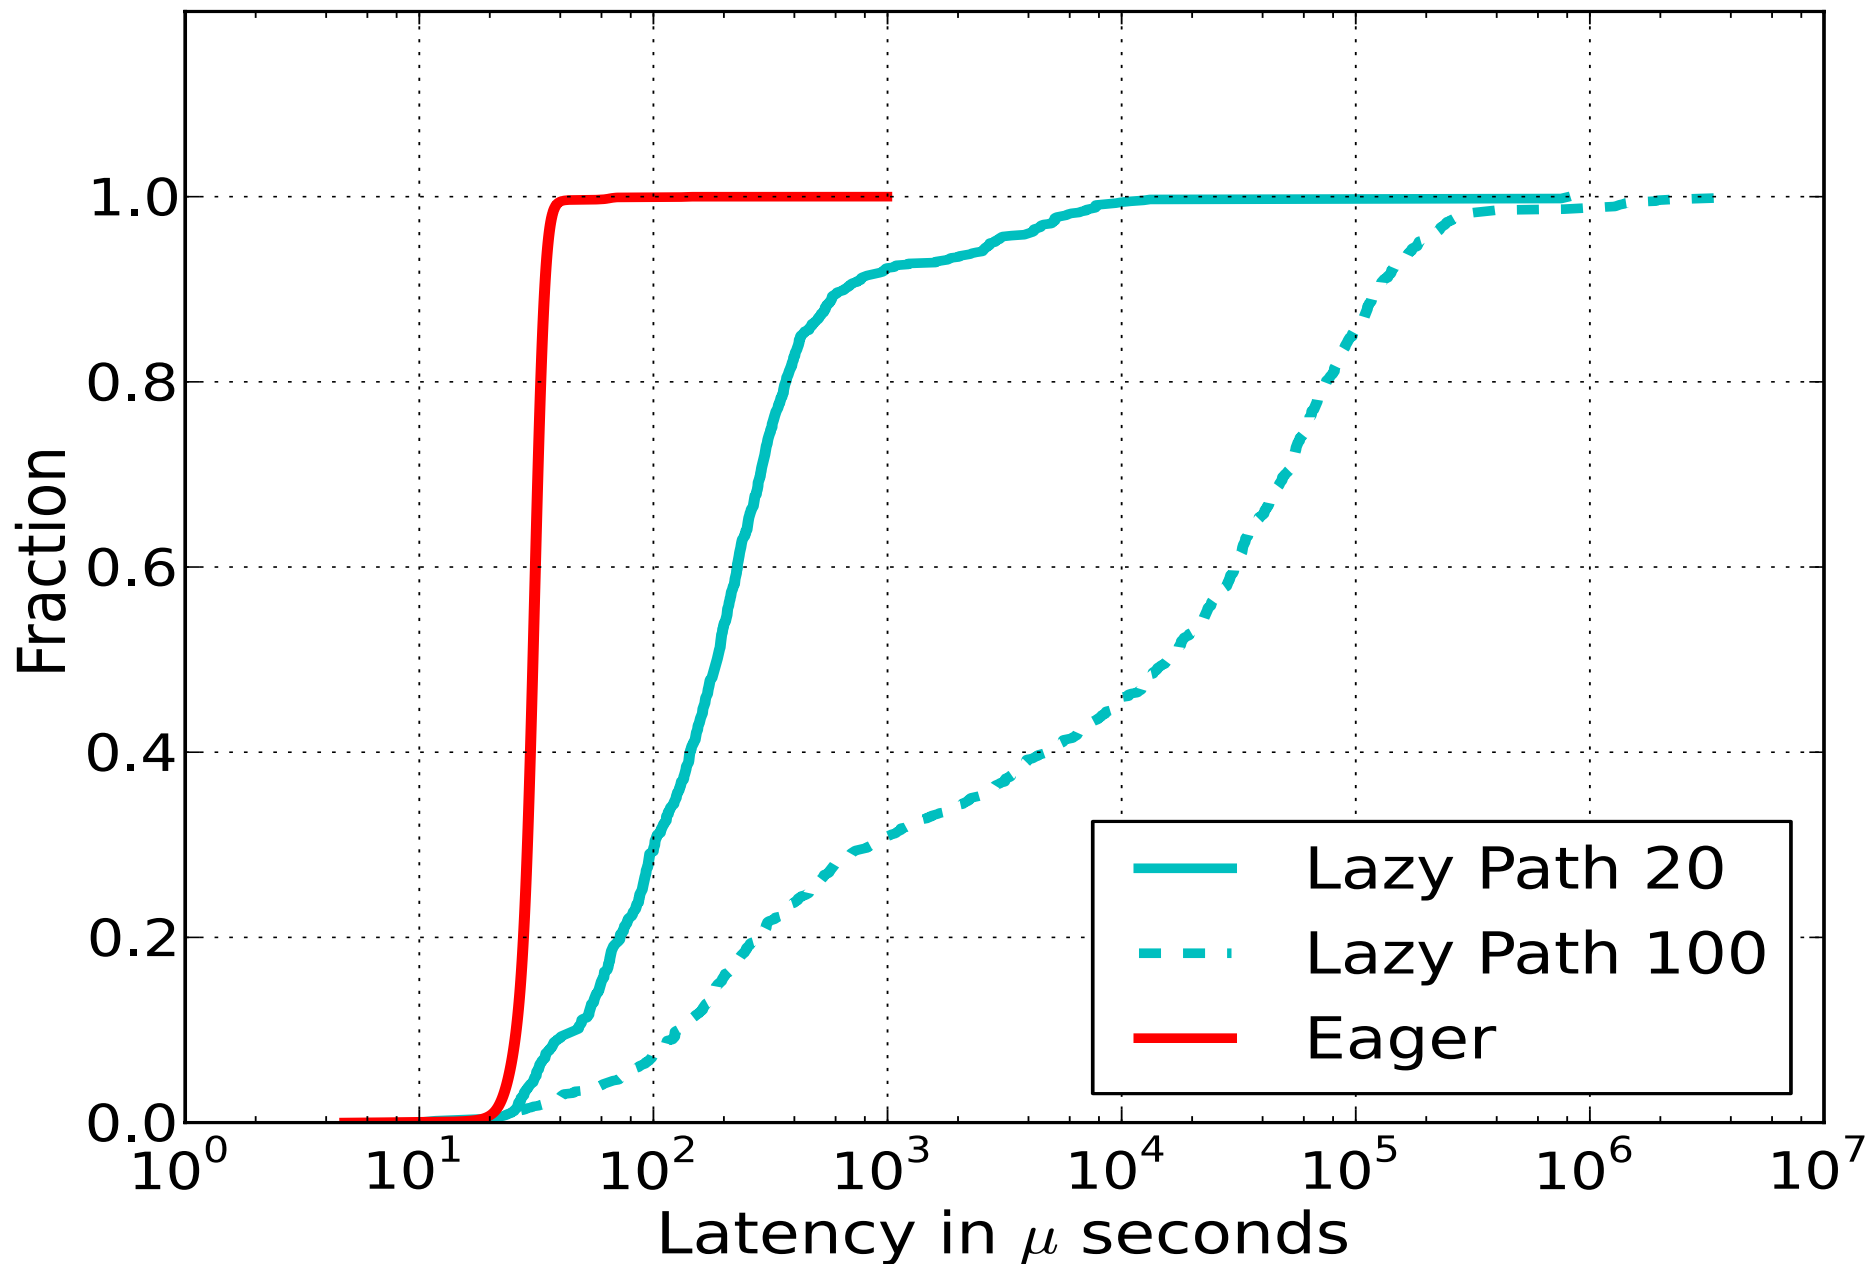
- “Cliquesy” workload
- Conflicting transactions conflict on several records



Cache-Locality



External Read Latency CDF

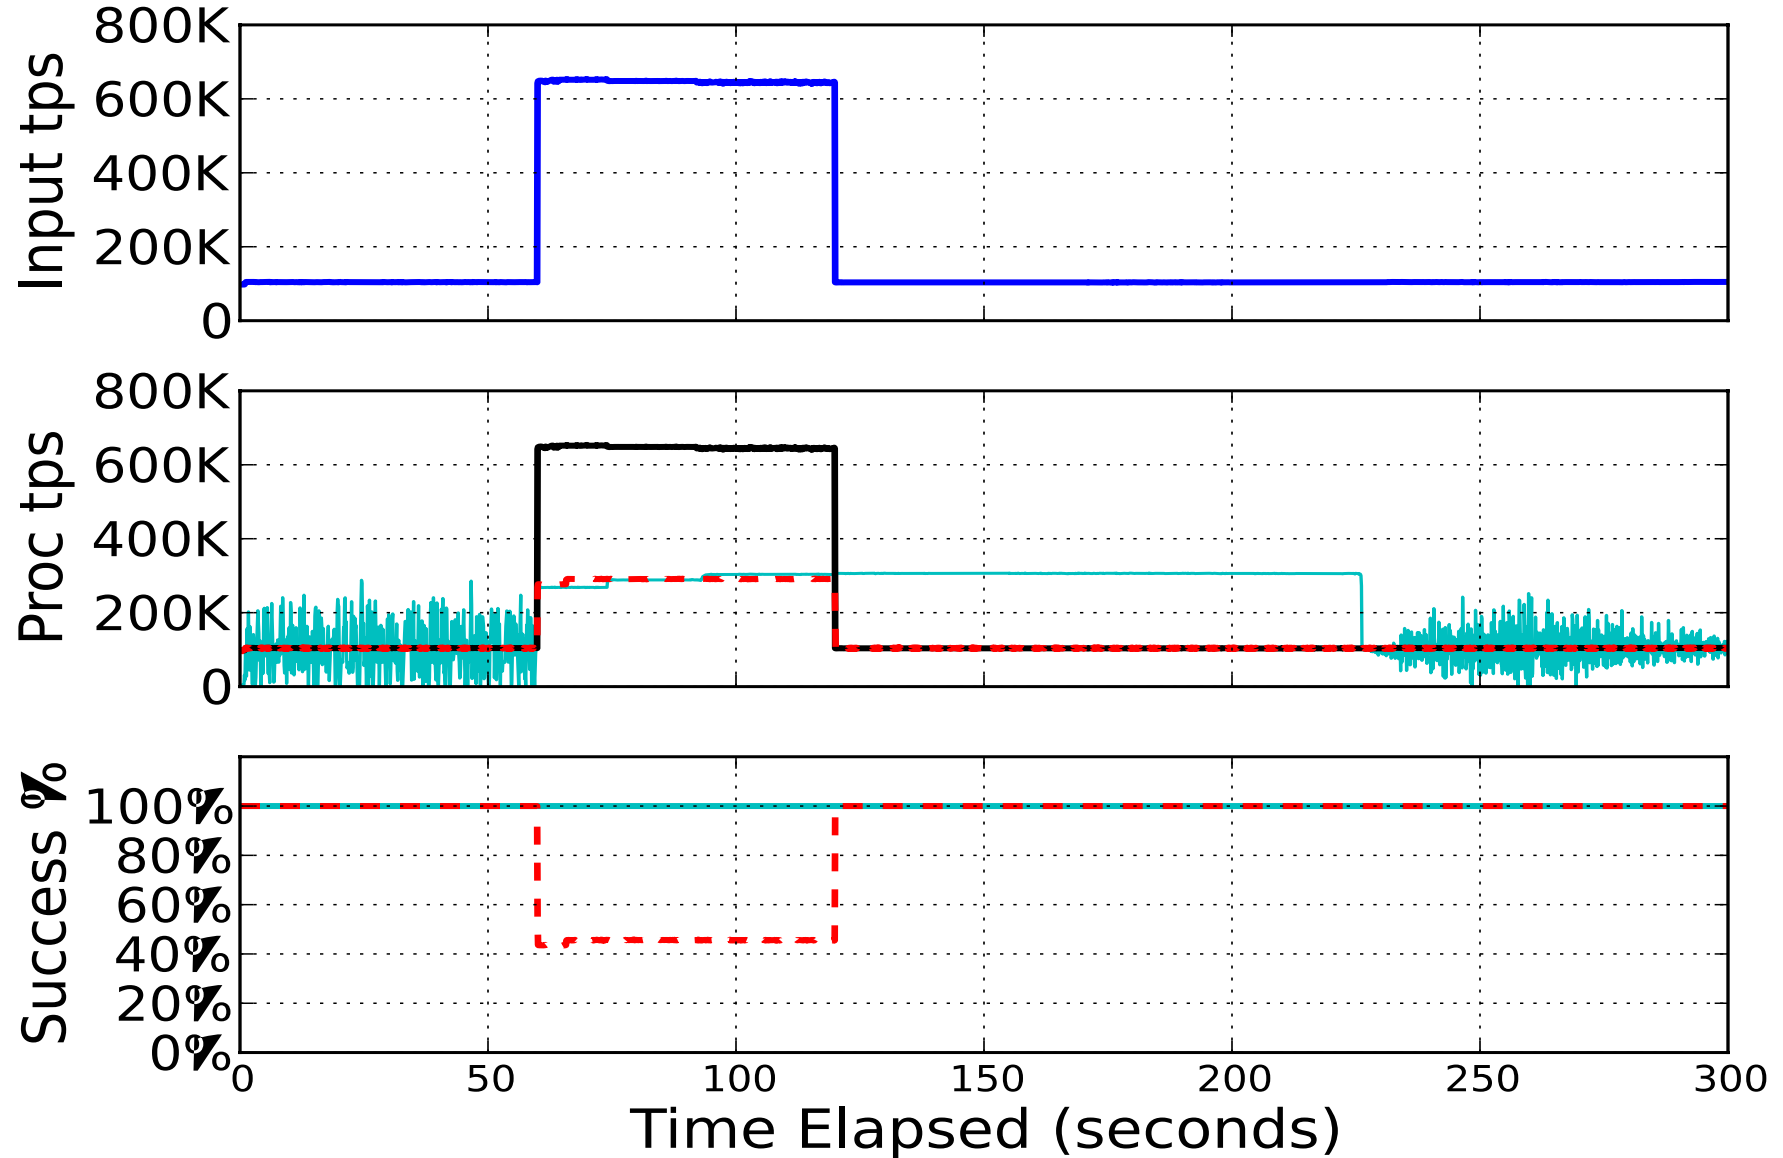


Temporal Load Balancing

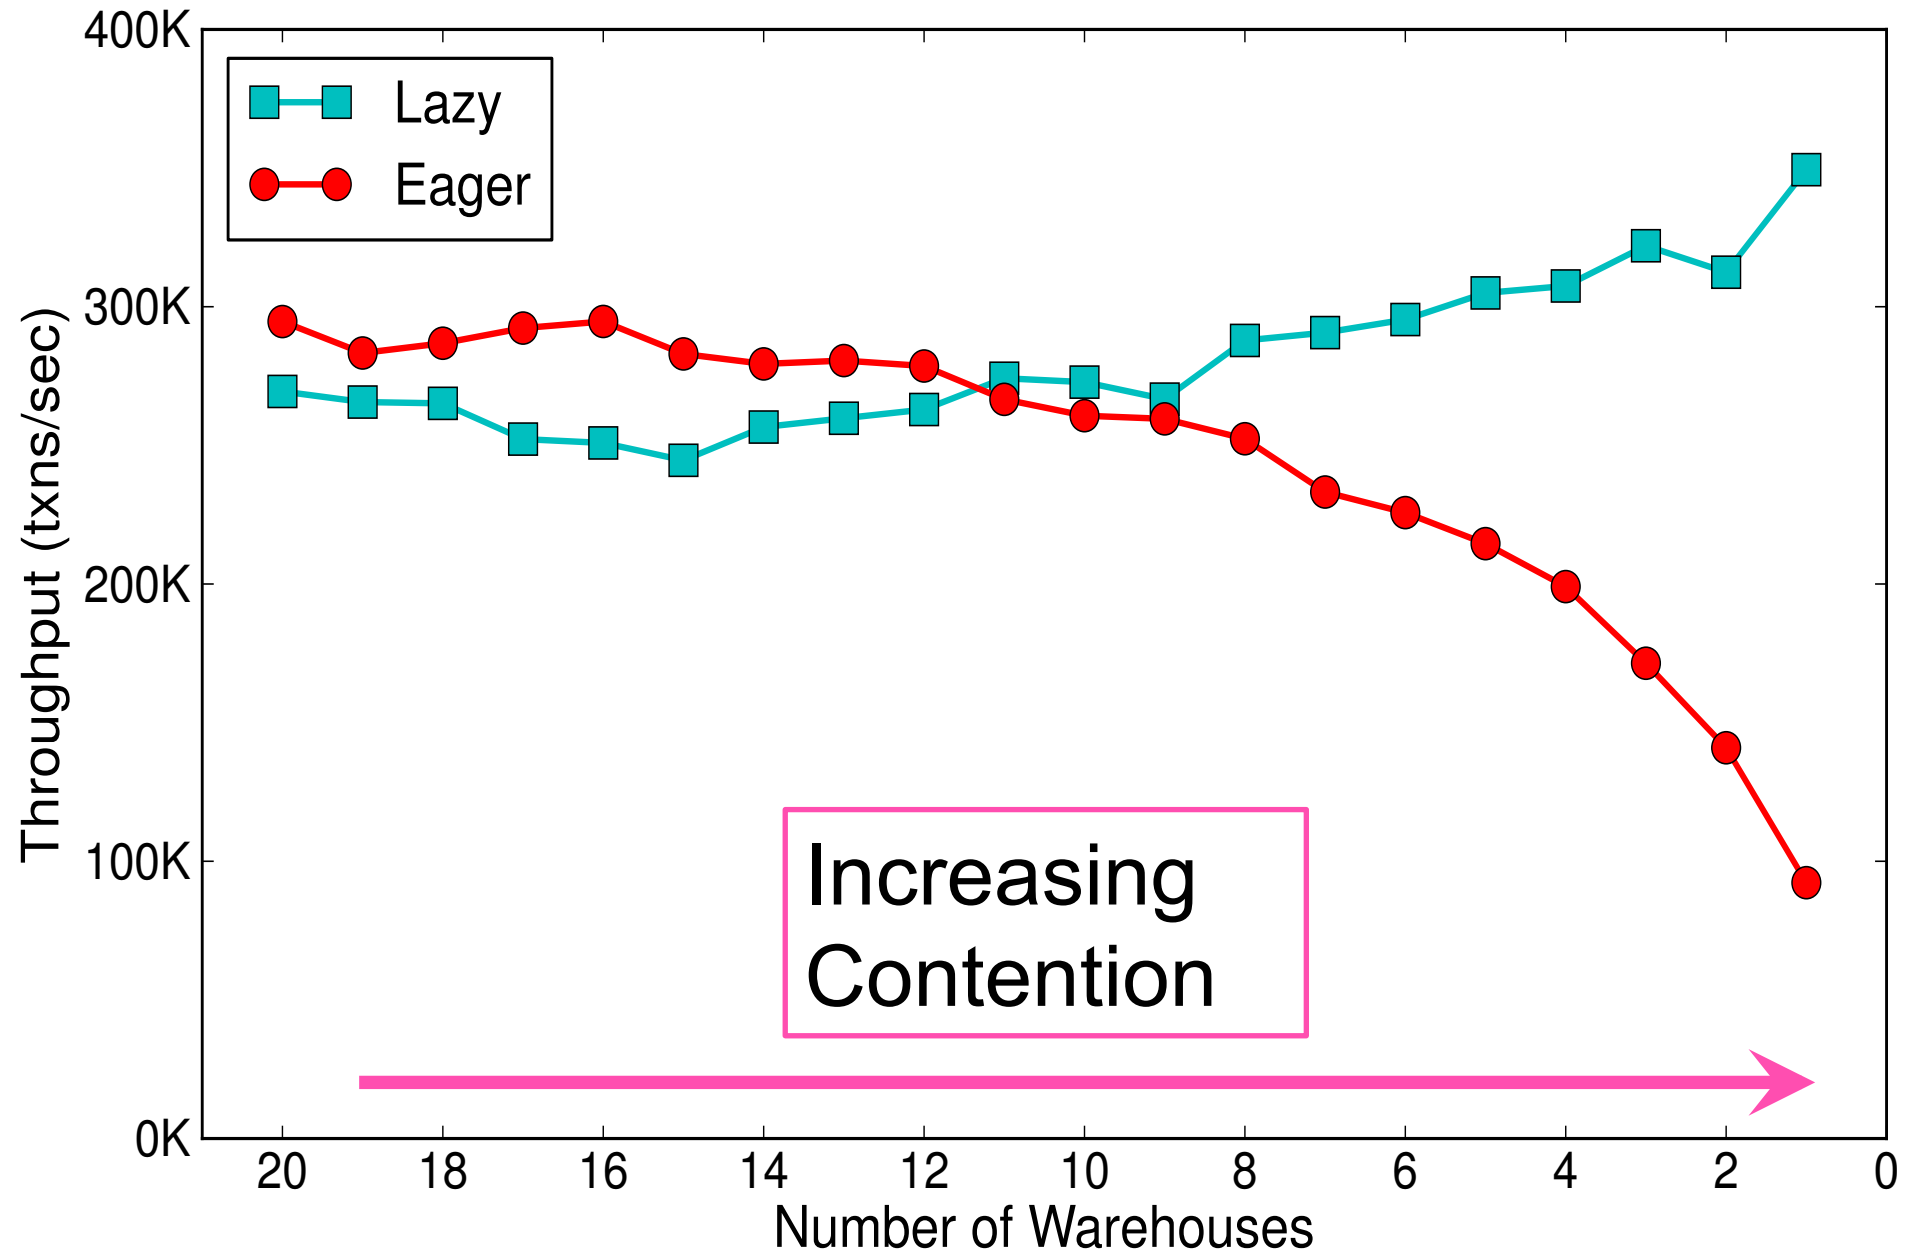
- Lazy and eager systems have comparable throughput

Temporal Load Balancing

— Input — Lazy - - Eager — Commit Decision



TPC-C Throughput



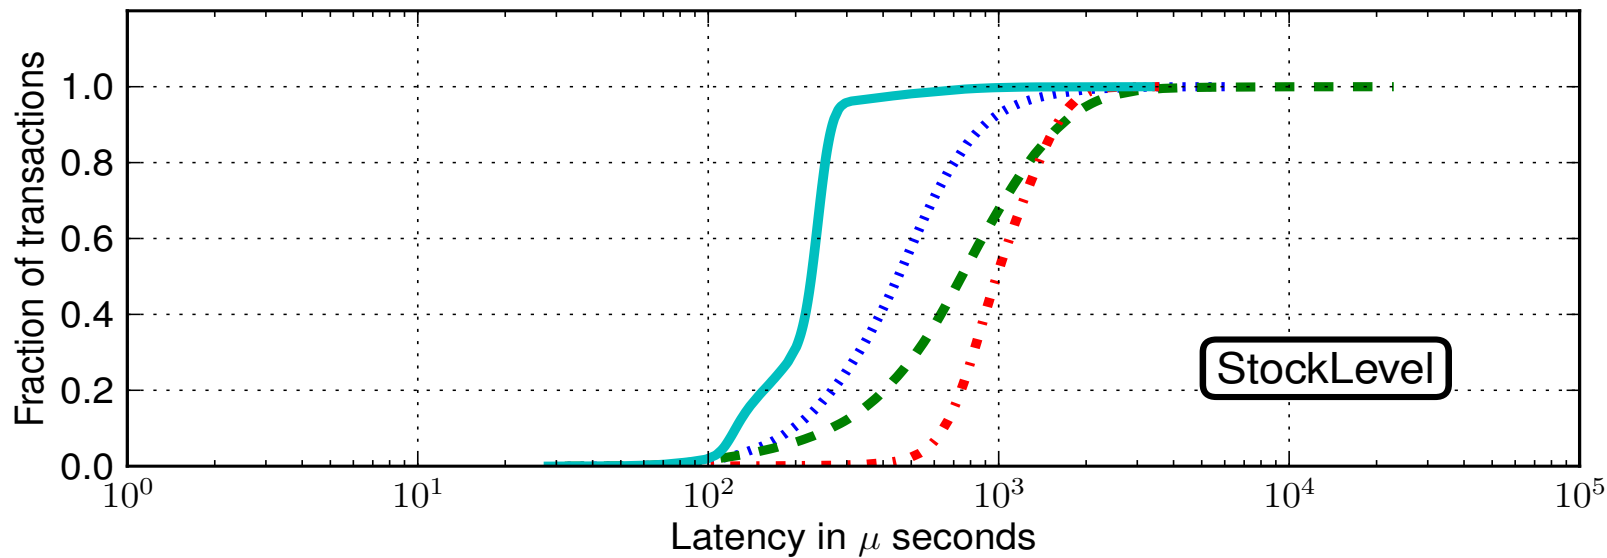

Conclusions

- Lazily evaluating txns has its benefits
 - Cache Locality
 - Temporal Load Balancing
 - Contention Reduction

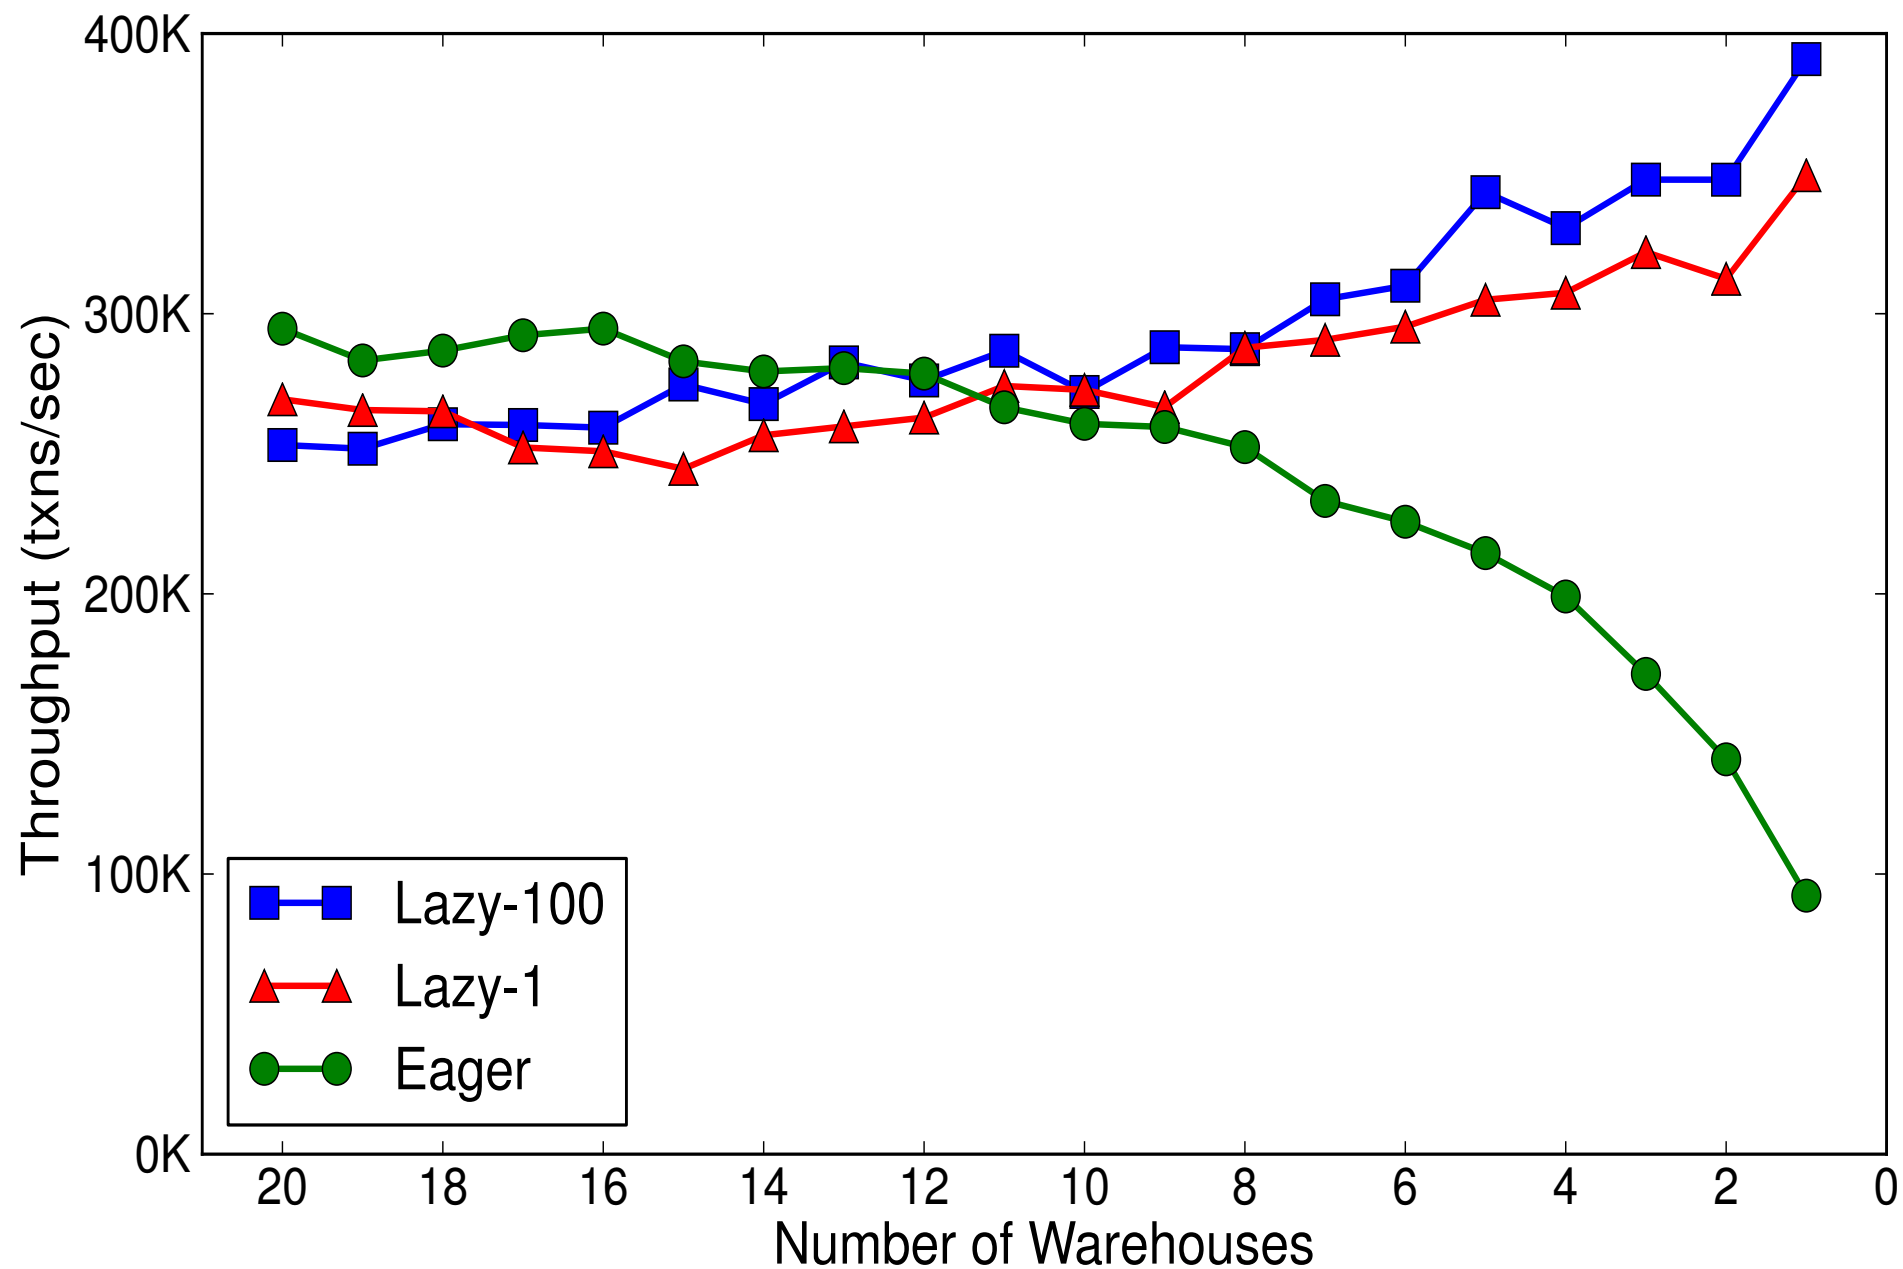
To see the rest of our
conclusions, read the
paper

TPC-C Latency

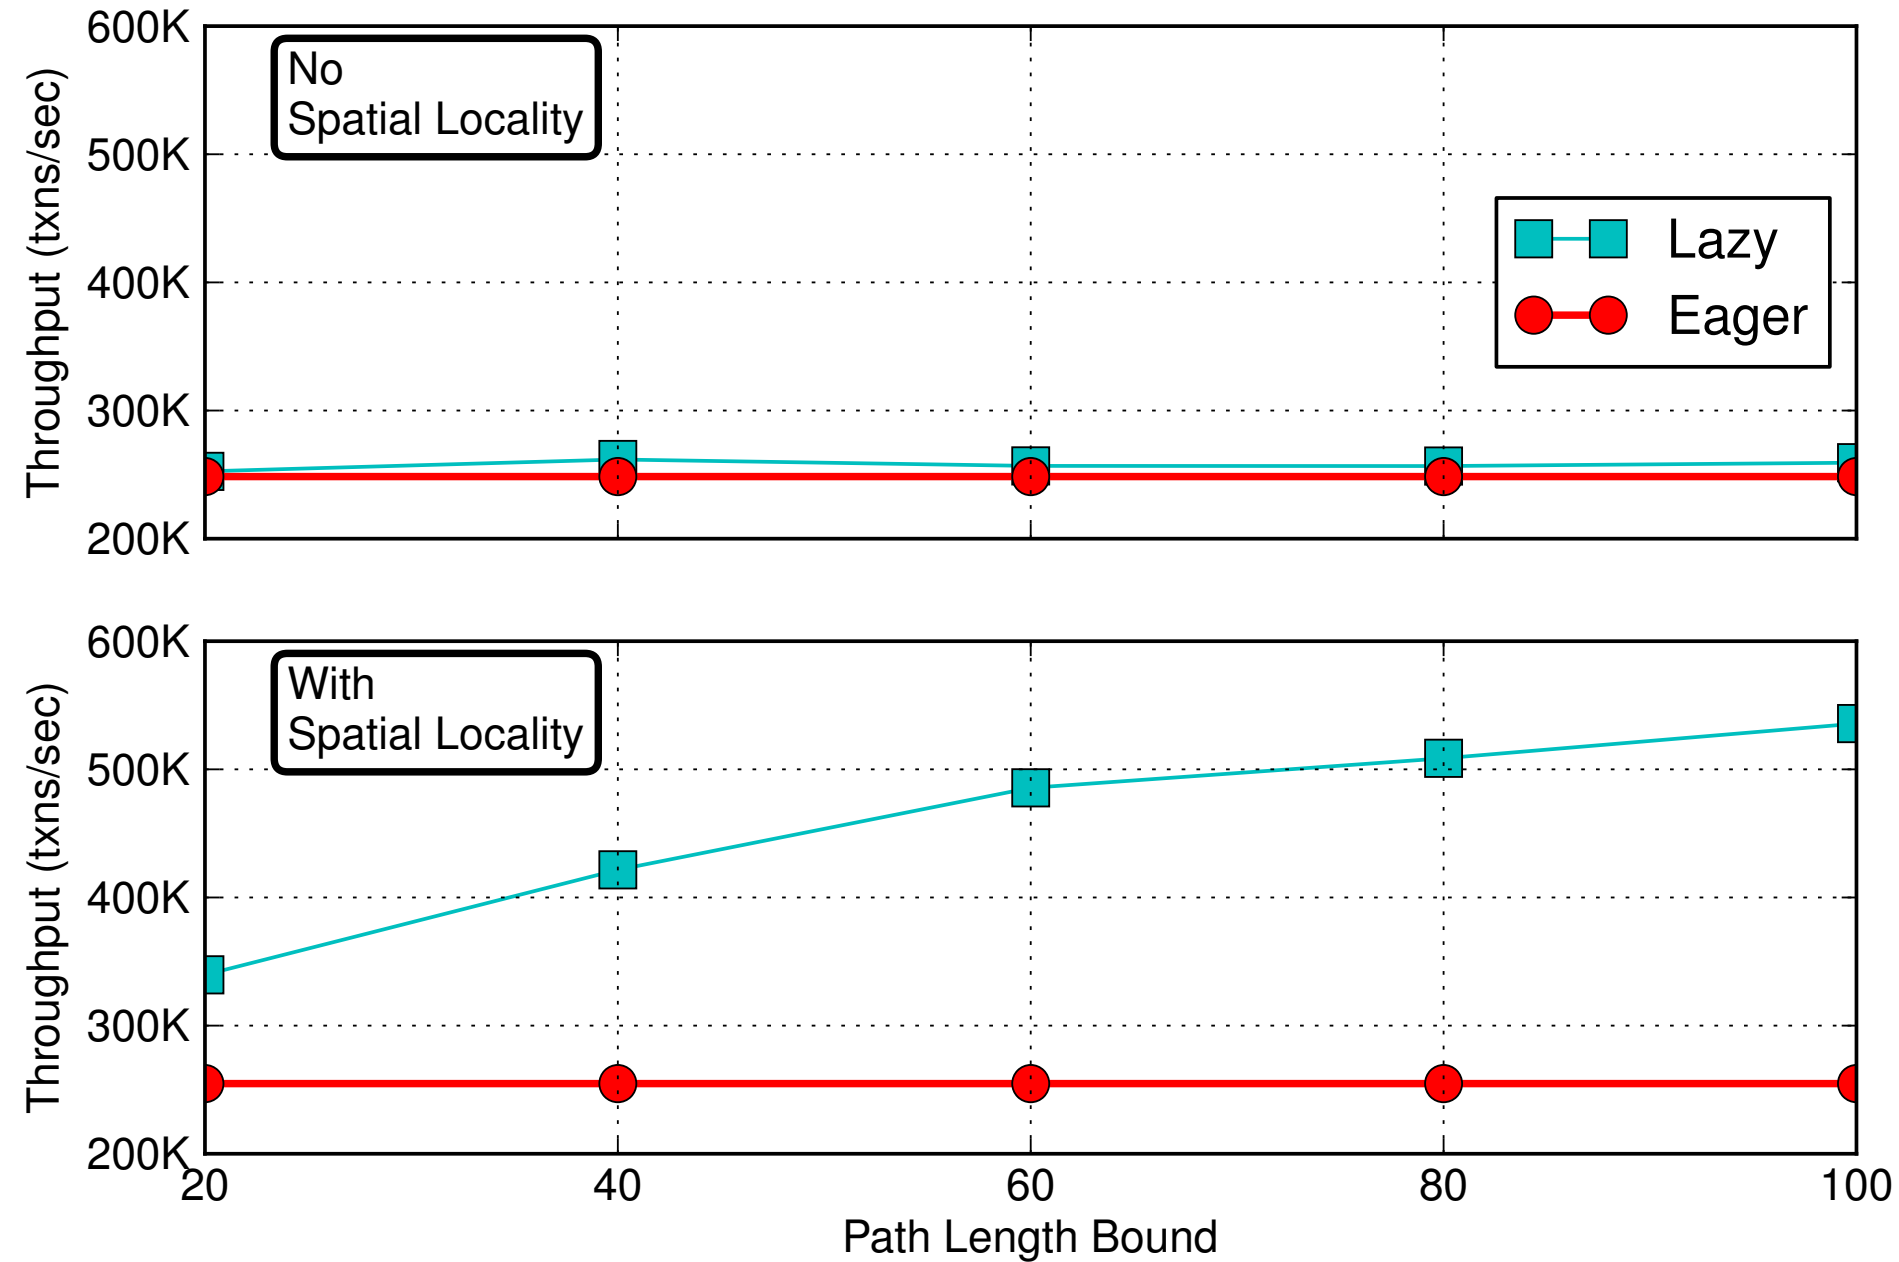
..... Lazy, 1 W - - - Lazy, 20 W . . . Eager, 1 W — Eager, 20 W



TPC-C Throughput



Microbenchmark Throughput



Microbenchmark Latency

Lazy Path 20 Lazy Path 100 Eager
Lazy Path 60

